

SQL Unit 3

Joins

Ryan Nixon

Joining Tables

- Allows for data to be pulled from multiple tables at once
- Connects the tables using their relationships
- Still produces a single resultset

The Importance of Aliases

- When working with two tables names can get very confusing. It's advisable to use aliases often to keep things understandable
- `SELECT first AS firstname
FROM person AS student`


Ambiguity

- When joining two tables that have columns with the same name, an ambiguity takes place
- The database doesn't know which column to work with and requires the column to be explicitly defined
- ```
SELECT p.first, p.last
FROM person AS p
WHERE p.age IS NOT NULL
```

# Joins Illustrated

| id | student_id | assn_id | points |
|----|------------|---------|--------|
| 1  | 1          | 1       | 5.00   |
| 2  | 6          | 3       | 27.00  |
| 3  | 6          | 1       | 9.00   |
| 4  | 5          | 2       | 13.50  |
| 5  | 3          | 3       | 25.00  |

| id | first       | last     | age |
|----|-------------|----------|-----|
| 1  | Jacob       | Williams | 37  |
| 2  | Emma        | Emmerich | 25  |
| 3  | Nathan      | Drake    | 32  |
| 4  | Albert      | Wesker   | 41  |
| 5  | Alexandra   | Vance    | 29  |
| 6  | Christopher | Garden   | 46  |



| id | student_id | assn_id | points | first       | last     | age |
|----|------------|---------|--------|-------------|----------|-----|
| 1  | 1          | 1       | 5.00   | Jacob       | Williams | 37  |
| 2  | 6          | 3       | 27.00  | Christopher | Garden   | 46  |
| 3  | 6          | 1       | 9.00   | Christopher | Garden   | 46  |
| 4  | 5          | 2       | 13.50  | Alexandra   | Vance    | 29  |
| 5  | 3          | 3       | 25.00  | Nathan      | Drake    | 32  |

# Unit 2 Assn

- Thinking back to Unit 2, there were many questions that required results from one query to be used in the second query
- Joins allow this to occur using a single query

# Two Join Syntaxes

- `SELECT *`  
`FROM person p, person_student s`  
`WHERE p.id = s.person_id`
- `SELECT *`  
`FROM person p`  
`INNER JOIN person_student s`  
`ON (p.id = s.person_id)`

# Two Join Syntaxes

- We will be using the second syntax for 2 reasons:
- The first syntax is typically only used for INNER JOINS. Additional types must use the second syntax
- The second syntax can handle much more complex queries without losing readability



# Inner Joins

- `INNER JOIN model`  
`ON (model.make_id = make.id)`
- Inner joins are the most common join
- Connects two tables with a boolean condition, just like `WHERE`
- Any connection that fails (i.e. there are no records on the “right” that exist for the “left”) do not show up in the final resultset

# Left & Right

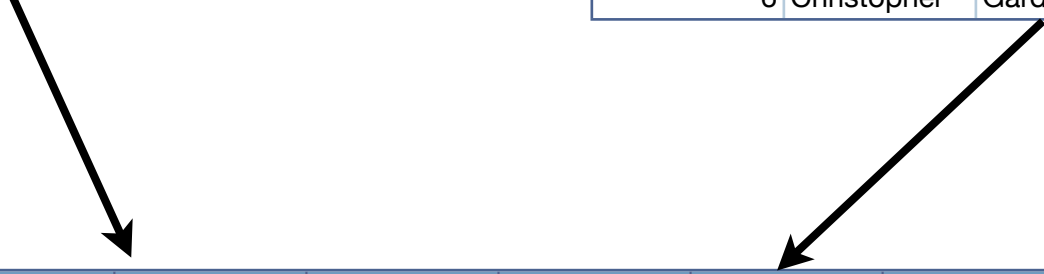
- The directions on a join correspond to the location of the table names
- `FROM person AS Left`  
`INNER JOIN person_student AS Right`
- Direction is extremely important. The left is the “original” data whereas the right is new data being added

# Inner Joins

```
SELECT *
FROM grade
INNER JOIN person
 ON person.id = grade.student_id
```

| id | student_id | assn_id | points |
|----|------------|---------|--------|
| 1  | 1          | 1       | 5.00   |
| 2  | 6          | 3       | 27.00  |
| 3  | 6          | 1       | 9.00   |
| 4  | 5          | 2       | 13.50  |
| 5  | 3          | 3       | 25.00  |

| id | first       | last     | age |
|----|-------------|----------|-----|
| 1  | Jacob       | Williams | 37  |
| 2  | Emma        | Emmerich | 25  |
| 3  | Nathan      | Drake    | 32  |
| 4  | Albert      | Wesker   | 41  |
| 5  | Alexandra   | Vance    | 29  |
| 6  | Christopher | Garden   | 46  |




| id | student_id | assn_id | points | first       | last     | age |
|----|------------|---------|--------|-------------|----------|-----|
| 1  | 1          | 1       | 5.00   | Jacob       | Williams | 37  |
| 2  | 6          | 3       | 27.00  | Christopher | Garden   | 46  |
| 3  | 6          | 1       | 9.00   | Christopher | Garden   | 46  |
| 4  | 5          | 2       | 13.50  | Alexandra   | Vance    | 29  |
| 5  | 3          | 3       | 25.00  | Nathan      | Drake    | 32  |

# Inner Joins


```
SELECT person.*, grade.*
FROM person
INNER JOIN grade
ON person.id = grade.student_id
```

| id | first       | last     | age |
|----|-------------|----------|-----|
| 1  | Jacob       | Williams | 37  |
| 2  | Emma        | Emmerich | 25  |
| 3  | Nathan      | Drake    | 32  |
| 4  | Albert      | Wesker   | 41  |
| 5  | Alexandra   | Vance    | 29  |
| 6  | Christopher | Garden   | 46  |

| id | student_id | assn_id | points |
|----|------------|---------|--------|
| 1  | 1          | 1       | 5.00   |
| 2  | 6          | 3       | 27.00  |
| 3  | 6          | 1       | 9.00   |
| 4  | 5          | 2       | 13.50  |
| 5  | 3          | 3       | 25.00  |



| id | first       | last     | age | student_id | assn_id | points |
|----|-------------|----------|-----|------------|---------|--------|
| 1  | Jacob       | Williams | 37  | 1          | 1       | 5.00   |
| 2  | Emma        | Emmerich | 25  |            |         |        |
| 3  | Nathan      | Drake    | 32  | 3          | 3       | 25.00  |
| 4  | Albert      | Wesker   | 41  |            |         |        |
| 5  | Alexandra   | Vance    | 29  | 5          | 2       | 13.50  |
| 6  | Christopher | Garden   | 46  | 6          | 3       | 27.00  |
| 6  | Christopher | Garden   | 46  | 6          | 1       | 9.00   |



| id | first       | last     | age | student_id | assn_id | points |
|----|-------------|----------|-----|------------|---------|--------|
| 1  | Jacob       | Williams | 37  | 1          | 1       | 5.00   |
| 3  | Nathan      | Drake    | 32  | 3          | 3       | 25.00  |
| 5  | Alexandra   | Vance    | 29  | 5          | 2       | 13.50  |
| 6  | Christopher | Garden   | 46  | 6          | 3       | 27.00  |
| 6  | Christopher | Garden   | 46  | 6          | 1       | 9.00   |

# Regarding Duplicates

- Note the duplication of Christopher Garden
- While a failed connection (NULL) will remove rows from the resultset, multiple connections will cause duplicates
- The duplication's purpose is to keep the ability to reference a single row without depending on other rows

# Other Join Names

- Equi-Join - The only condition connecting the tables is the equality of their keys
- Natural Join - If the columns performing the connections are equal, merge them into a single column
- Cross Join - No conditions are specified; this causes the resultset to be the cartesian product of the two tables

# Join Conditions


- Equi-Joins are not required. Additional conditions may be supplied to further limit results
- In rare cases this improves performance over using filters in the WHERE clause
- In some products, the WHERE is applied to the final resultset after all joins have taken place
- JOINS apply their filters as they build the results, keeping the final resultset small

# Join Conditions

```
SELECT *
FROM grade
INNER JOIN person
 ON (person.id = grade.student_id AND person.age > 32)
```

| id | student_id | assn_id | points |
|----|------------|---------|--------|
| 1  | 1          | 1       | 5.00   |
| 2  | 6          | 3       | 27.00  |
| 3  | 6          | 1       | 9.00   |
| 4  | 5          | 2       | 13.50  |
| 5  | 3          | 3       | 25.00  |

| id | first       | last     | age |
|----|-------------|----------|-----|
| 1  | Jacob       | Williams | 37  |
| 2  | Emma        | Emmerich | 25  |
| 3  | Nathan      | Drake    | 32  |
| 4  | Albert      | Wesker   | 41  |
| 5  | Alexandra   | Vance    | 29  |
| 6  | Christopher | Garden   | 46  |



| id | student_id | assn_id | points | first       | last     | age |
|----|------------|---------|--------|-------------|----------|-----|
| 1  | 1          | 1       | 5.00   | Jacob       | Williams | 37  |
| 2  | 6          | 3       | 27.00  | Christopher | Garden   | 46  |
| 3  | 6          | 1       | 9.00   | Christopher | Garden   | 46  |
| 4  | 5          | 2       | 13.50  | Alexandra   | Vance    | 29  |
| 5  | 3          | 3       | 25.00  | Nathan      | Drake    | 32  |



# Multiple Joins

- There is no limit to the number of joins that can be specified. Any join will work off of the resultset from the prior join. This allows columns from the Right table to be used in the next join
- This is perfect for Many-to-Many relationships

# Multiple Joins

```
SELECT first, last, name AS assignment, points
FROM person
INNER JOIN grade ON (grade.student_id = person.id)
INNER JOIN assn ON (assn.id = grade.assn_id)
ORDER BY assn.due, assn.name, person.last
```

| id | first       | last     | age |
|----|-------------|----------|-----|
| 1  | Jacob       | Williams | 37  |
| 2  | Emma        | Emmerich | 25  |
| 3  | Nathan      | Drake    | 32  |
| 4  | Albert      | Wesker   | 41  |
| 5  | Alexandra   | Vance    | 29  |
| 6  | Christopher | Garden   | 46  |

| id | student_id | assn_id | points |
|----|------------|---------|--------|
| 1  | 1          | 1       | 5.00   |
| 2  | 6          | 3       | 27.00  |
| 3  | 6          | 1       | 9.00   |
| 4  | 5          | 2       | 13.50  |
| 5  | 3          | 3       | 25.00  |

| id | name   | due        |
|----|--------|------------|
| 1  | Unit 1 | 09/24/2012 |
| 2  | Unit 2 | 09/24/2012 |
| 3  | Unit 3 | 09/24/2012 |
| 4  | Test 1 | 09/24/2012 |

| first       | last     | assignment | points |
|-------------|----------|------------|--------|
| Christopher | Garden   | Unit 1     | 9.00   |
| Jacob       | Williams | Unit 1     | 5.00   |
| Alexandra   | Vance    | Unit 2     | 13.50  |
| Nathan      | Drake    | Unit 3     | 25.00  |
| Christopher | Garden   | Unit 3     | 27.00  |

# Self-Joining

- Tables can also be joined with themselves. This adds further importance to the use of aliases
- A perfect example of this is the parent-child relationship, where both sets of data are considered people

# Self-Joining

```
SELECT parent.first AS parent_first, parent.last AS parent_last,
 child.first AS child_first, child.last AS child_last
FROM person AS parent
INNER JOIN person AS child
 ON (child.father_id = parent.id
 OR child.mother_id = parent.id)
```

| id | mother_id | father_id | first  | last     | age  |
|----|-----------|-----------|--------|----------|------|
| 1  | 4         | 3         | Todd   | Flanders | 8    |
| 2  | 6         | 5         | Stewie | Griffin  | 1    |
| 3  | NULL      | NULL      | Ned    | Flanders | 47   |
| 4  | NULL      | NULL      | Maude  | Flanders | NULL |
| 5  | NULL      | NULL      | Peter  | Griffin  | 42   |
| 6  | NULL      | NULL      | Lois   | Griffin  | 39   |



| parent_first | parent_last | child_first | child_last |
|--------------|-------------|-------------|------------|
| Ned          | Flanders    | Todd        | Flanders   |
| Maude        | Flanders    | Todd        | Flanders   |
| Peter        | Griffin     | Stewie      | Griffin    |
| Lois         | Griffin     | Stewie      | Griffin    |

# Self-Joining

```
SELECT child.first AS child_first, child.last AS child_last,
 mother.first AS mother_first, mother.last AS mother_last,
 father.first AS father_first, father.last AS father_last
FROM person AS child
INNER JOIN person AS mother ON (mother.id = child.mother_id)
INNER JOIN person AS father ON (father.id = child.father_id)
```

| id | mother_id | father_id | first  | last     | age  |
|----|-----------|-----------|--------|----------|------|
| 1  | 4         | 3         | Todd   | Flanders | 8    |
| 2  | 6         | 5         | Stewie | Griffin  | 1    |
| 3  | NULL      | NULL      | Ned    | Flanders | 47   |
| 4  | NULL      | NULL      | Maude  | Flanders | NULL |
| 5  | NULL      | NULL      | Peter  | Griffin  | 42   |
| 6  | NULL      | NULL      | Lois   | Griffin  | 39   |



| child_first | child_last | mother_first | mother_last | father_first | father_last |
|-------------|------------|--------------|-------------|--------------|-------------|
| Todd        | Flanders   | Maude        | Flanders    | Ned          | Flanders    |
| Stewie      | Griffin    | Lois         | Griffin     | Peter        | Griffin     |

# Outer Joins

- Exactly like Inner Joins, except in the case where connections fail. When a NULL is encountered the empty data will be left in the resultset
- Outer joins also need their direction specified. They use LEFT JOIN or RIGHT JOIN for this purpose

# Outer Joins

- With a LEFT JOIN, any data on the left side that doesn't have a corresponding row will stay in the resultset, i.e. NULLs on the right will be preserved
- With a RIGHT JOIN, data on the right side will be present in the resultset, regardless of the connection. This causes NULLs to appear in the original resultset

# Compatibility

- For the purposes of this class, RIGHT JOINS will not be used as they are not fully supported by all databases



# Outer Joins

```
SELECT *
FROM person AS child
LEFT JOIN person AS mother ON (mother.id = child.mother_id)
LEFT JOIN person AS father ON (father.id = child.father_id)
```

| id | mother_id | father_id | first  | last     | age  |
|----|-----------|-----------|--------|----------|------|
| 1  | 4         | 3         | Todd   | Flanders | 8    |
| 2  | 6         | 5         | Stewie | Griffin  | 1    |
| 3  | NULL      | NULL      | Ned    | Flanders | 47   |
| 4  | NULL      | NULL      | Maude  | Flanders | NULL |
| 5  | NULL      | NULL      | Peter  | Griffin  | 42   |
| 6  | NULL      | NULL      | Lois   | Griffin  | 39   |



| first  | last     | first | last     | first | last     |
|--------|----------|-------|----------|-------|----------|
| Todd   | Flanders | Maude | Flanders | Ned   | Flanders |
| Stewie | Griffin  | Lois  | Griffin  | Peter | Griffin  |
| Ned    | Flanders | NULL  | NULL     | NULL  | NULL     |
| Maude  | Flanders | NULL  | NULL     | NULL  | NULL     |
| Peter  | Griffin  | NULL  | NULL     | NULL  | NULL     |
| Lois   | Griffin  | NULL  | NULL     | NULL  | NULL     |

# Outer Joins

- Because outer joins never remove data from the resultset, they are excellent for statistics and data aggregation
- “All students and classes they are taking, *if any*”
- “All sales that these car models have had, *if any*”
- In these cases you would want to still have the original source intact in your data to track any non-participants or low-sellers

# Reminders

- Assignment 3 up tonight. Due 9/24
- Lab time Wednesday in SSB 172
- Test Monday in this classroom on 9/24
- Additional office hours on Sunday in SSB 172, 2pm-4pm