

SQL Unit 7

Set Operations

Ryan Nixon

Another Way To Select

- Recall that Joins allow us to combine the rows from two tables into a single resultset
- WHERE can filter those results
- SELECT will display only the columns we choose
- It turns out there is one more way to specify how your data is handled


Set Operations

- Set operations -- Working with result**sets** as if they were collections of items, or tuples
- Using set theory, we can combine or intersect sets to make substantial changes to the resultset

Back to WHERE


- Remember AND and OR in the WHERE clause?
 - AND will require both the condition on the left and the condition on the right to be true
 - OR will require one condition on the left or right to be true, or both
- Sets work similar to this but at a higher level; instead of working with columns, they work with whole rows

Set Notation




- You may already be familiar with set notation. The symbols used in those directly correlate to their SQL implementations:
 - $A \cup B$ - A UNION B (SQL OR)
 - $A \cap B$ - A INTERSECTION B (SQL AND)

Set Notation



- $A \cup B$ - A UNION B (SQL OR)
- Return the Union of query A and query B. This will include all rows that occur in either resultset or both of them

Set Notation




- $A \cap B$ - A INTERSECTION B (SQL AND)
- Return the intersection of query A and query B. This will include only rows that occur in *both* resultsets

Set Notation

- In practice, the set operations will also remove duplicate rows that are encountered. This is similar to applying **DISTINCT** to the query

Unions



- ```
SELECT *
FROM person
WHERE last = 'Smith'
UNION
SELECT *
FROM person
WHERE last = 'Jones'
```

# Unions

A

| person   |         |       |           |            |
|----------|---------|-------|-----------|------------|
| id       | first   | last  | gender_id | address_id |
| 30663453 | Lisa    | Smith | F         | 738        |
| 30729595 | Allen   | Smith | M         | 185        |
| 30782165 | Norbert | Smith | M         | 735        |

U

| person   |          |       |           |            |
|----------|----------|-------|-----------|------------|
| id       | first    | last  | gender_id | address_id |
| 30600610 | Joyce    | Jones | F         | 520        |
| 30717694 | Steve    | Jones | M         | 260        |
| 30799275 | Ken      | Jones | M         | 559        |
| 30840089 | Benjamin | Jones | M         | 100        |

B



| person   |          |       |           |            |
|----------|----------|-------|-----------|------------|
| id       | first    | last  | gender_id | address_id |
| 30600610 | Joyce    | Jones | F         | 520        |
| 30663453 | Lisa     | Smith | F         | 738        |
| 30717694 | Steve    | Jones | M         | 260        |
| 30729595 | Allen    | Smith | M         | 185        |
| 30782165 | Norbert  | Smith | M         | 735        |
| 30799275 | Ken      | Jones | M         | 559        |
| 30840089 | Benjamin | Jones | M         | 100        |

# Unions

- Simply putting UNION between the two queries will trigger the set operation
- Take a look at the query again. Doesn't this behavior look familiar? Is there another way to do it?

# Unions


- ```
SELECT *  
FROM person  
WHERE last = 'Smith' OR last = 'Jones'
```
- This would produce the same resultset. In this way UNION acts exactly like the OR.
- UNION however is very beneficial when more complex queries are in use

UNION ALL

- There is another modifier that can be used, UNION ALL, which will keep any duplicates around.
- This may be useful if the results are going to be aggregated in some way

Intersections

- ```
SELECT *
FROM class
WHERE credits > 1
LIMIT 5
INTERSECT
SELECT *
FROM class
WHERE credits < 3
LIMIT 5
```



# Intersections

A

| class |                                         |        |         |
|-------|-----------------------------------------|--------|---------|
| id    | name                                    | course | credits |
| 70004 | Principles of Financial Accounting I    | A101   | 3       |
| 70006 | Principles of Financial Accounting II   | A102   | 3       |
| 70007 | Bookkeeping for Business I              | A120   | 3       |
| 70014 | Principles of Financial Accounting      | A201   | 3       |
| 70017 | Principles of Managerial Accounting     | A202   | 3       |
| 70020 | Introduction to Computerized Accounting | A222   | 3       |

$\cap$

| class |                                              |        |         |
|-------|----------------------------------------------|--------|---------|
| id    | name                                         | course | credits |
| 70116 | Foundations of the United States Air Force I | A101   | 1       |
| 70117 | US Air Force Leadership Laboratory           | A150   | 1       |
| 70118 | Evolution of Air and Space Power I           | A201   | 2       |
| 70128 | Aircraft Ground Operations and Safety        | A170   | 1       |
| 70132 | Fundamentals of Aircraft Electronics Lab     | A174L  | 2       |
| 70133 | Drawing and Precision Measurement            | A175   | 2       |

B

| class |                                          |        |         |
|-------|------------------------------------------|--------|---------|
| id    | name                                     | course | credits |
| 70118 | Evolution of Air and Space Power I       | A201   | 2       |
| 70132 | Fundamentals of Aircraft Electronics Lab | A174L  | 2       |
| 70133 | Drawing and Precision Measurement        | A175   | 2       |

# Intersections

- Intersections have a very similar syntax to Unions, but a completely different behavior
- Again, take a look at the query. What does that look familiar to?



# Intersections

- `SELECT *`  
`FROM class`  
`WHERE credits > 1 AND credits < 3`
- Intersections are exactly like the AND operator when they work with resultsets
- Just like Unions, this only works well in complex queries

# A Compatibility Note

- For the sake of illustration, I just lied a bit
- The INTERSECT keyword does not work in many SQL implementations (although it does work in SQLite)
- But that's okay! After this short tangent you'll see something *BETTER!*

# Dealing With Columns

- Note that the previous examples all use \* and are from the same table. This is not a requirement
- As long as the types are similar and the number of columns are the same, you may interchange tables
- This is known as union compatibility

# Dealing With Columns

- Note that I said 'similar' when dealing with types. There is an implicit cast that occurs when dealing with columns of similar type
- $\text{CHAR}(10) + \text{CHAR}(50) = \text{CHAR}(50)$
- $\text{INT} + \text{FLOAT} = \text{FLOAT}$
- The cast will take the largest or most flexible field and use that in the resultset

# Dealing With Columns

- With columns, the counts must be the same. As this is unlikely to occur when tables are being compared, you can “pad” columns with default values until they match
- `person_student: person_id, NULL, out_of_state, NULL`
- `person_faculty: person_id, department_id, 0, bio`
- Note that this value doesn't have to be NULL

# Dealing With Columns

- ```
SELECT person_id, NULL, out_of_state, NULL
FROM person_student
LIMIT 5
UNION
SELECT person_id, department_id, 0, bio
FROM person_faculty
LIMIT 5
```

Dealing With Columns

A

person_student			
person_id	NULL	out_of_state	NULL
30676986		0	
30677070		0	
30679836		0	
30680792		1	
30681904		0	



U

person_faculty			
person_id	department_id	0	bio
30745137	7	0	
30745820	1	0	
30748325	30	0	
30752649	24	0	I like cats.
30756264	9	0	



B


person_id	NULL	out_of_state	NULL
30676986		0	
30677070		0	
30679836		0	
30680792		1	
30681904		0	
30745137	7	0	
30745820	1	0	
30748325	30	0	
30752649	24	0	I like cats.
30756264	9	0	

Dealing With Columns

- You might note that the header for the resultset had NULLs in it.
- It takes the headers from query A, so if you would like better headers you can add aliases to the NULL columns

Additional Operations

- Union and Intersect directly relate to WHERE operators, but there are 2 more that indirectly relate:
 - $A \subset B$ - A IN B
 - $A - B$ - A NOT IN B



- IN...Intersect...oh, I see what you did there.
- It's true, they perform the same actions, but in much different ways
- The IN operator, of the 4 in this unit, is by far the most common

IN


- ```
SELECT *
FROM person_student
WHERE person_id IN (
 SELECT person_id
 FROM person_faculty)
LIMIT 5
```

# IN

| person_student |              |
|----------------|--------------|
| person_id      | out_of_state |
| 30770668       | 0            |
| 30707731       | 0            |
| 30752649       | 0            |
| 30799694       | 0            |
| 30700976       | 0            |

⊂

| person_faculty |  |
|----------------|--|
| person_id      |  |
| 30748325       |  |
| 30752649       |  |
| 30756264       |  |
| 30759542       |  |
| 30759864       |  |



| person_id | out_of_state |
|-----------|--------------|
| 30752649  |              |

# IN

- Notice how the operation behaves the same, selecting all rows that appear in both, but the syntax is vastly different
- With IN you can intersect two resultsets based off of a common column. Almost like a join, but in this case no additional rows or columns are added

# A Familiar Example

- ```
SELECT *  
FROM person  
INNER JOIN class_registration reg  
    ON (reg.person_id = person.id  
        AND reg.role_id = 'S')  
INNER JOIN class_registration reg2  
    ON (reg2.person_id = person.id  
        AND reg2.role_id = 'F')
```

A Familiar Example

- ```
SELECT *
FROM person
INNER JOIN class_registration reg
 ON (reg.person_id = person.id)
WHERE role_id = 'S'
 AND person_id IN (
 SELECT person_id
 FROM class_registration
 WHERE role_id = 'F')
```

# IN


- Also notice that there is an entire query in those parentheses. This (surprise) is what an inline view or subquery looks like
- It defines another query whose resultset is used as a parameter for IN
- Even better: since the resultset is only one column, you can substitute the query for actual values



# IN

- `SELECT *`  
`FROM person`  
`WHERE last IN ( 'Smith', 'Jones' )`
- This query is exactly like the Union example. It filters all person rows down to those with the last name 'Smith' OR 'Jones'

# NOT IN



- You can apply the NOT operator to IN as well:
- ```
SELECT *  
FROM person  
WHERE last NOT IN ('Smith',  
  'Jones')
```
- This predictably will return all persons whose last name is NOT Smith or Jones

Reminders

- Assignment 7 up by Friday. Due 11/19
- Independent lab time Monday in SSB 172
- Lecture next Wednesday in classroom